

Secure Model Supply Chain for Enterprise AI on HPE: Attestation and Behavioral Fingerprinting of AI Models Before Deployment

[Author Name]

HPE [Business Unit], Bengaluru, India

firstname.lastname@hpe.com

Abstract

Enterprises deploying AI workloads on HPE GreenLake and HPE Ezmeral Container Platform routinely import models from HuggingFace, GitHub, and third-party vendors with no systematic integrity verification before deployment. This creates a critical supply chain risk: model weights may be backdoored or poisoned at source, and no enterprise ML serving stack enforces a trust gate prior to scheduling. We present SecureModelGate, a pre-deployment attestation pipeline that produces a cryptographically signed Model Attestation Token (MAT) binding a model's weight hash to its behavioral fingerprint on a canonical test corpus, alongside a signed Model Bill of Materials (MBOM). The MAT is enforced by a Kubernetes validating admission webhook on HPE Ezmeral, blocking any unattested model-serving pod from scheduling. Evaluated on 120 models from the NIST TrojAI benchmark, our system achieves 91.3% backdoor detection at 4.1% false positive rate with under 40 seconds attestation latency - the first pipeline to combine static analysis, behavioral fingerprinting, and MBOM generation into an enforceable deployment gate within HPE enterprise container infrastructure.

Problem statement

Enterprises running AI on HPE GreenLake, Ezmeral Container Platform, and HPC clusters regularly pull model artifacts from HuggingFace, GitHub, and vendors with no integrity verification. HuggingFace alone hosts over 500,000 repositories [1] without verifiable provenance, creating a supply chain attack surface more dangerous than software dependencies.

Backdoored models produce incorrect outputs only on specific trigger inputs, evading conventional testing [2, 3]. MITRE ATLAS [5] recognizes AI supply chain attacks, yet no enterprise ML stack - MLflow [4], KServe, or otherwise - enforces a mandatory pre-deployment trust gate. For HPE's regulated customers (banking, healthcare, defence), this gap carries significant liability.

Existing tools address only file integrity (Sigstore/cosign [6]) or offer detection algorithms without enforcement (Neural Cleanse [7], STRIP [8]). No tool combines behavioral attestation with an enforceable deployment gate on HPE infrastructure.

Our solution

We introduce SecureModelGate, a pre-deployment model attestation pipeline purpose-built for HPE Ezmeral and GreenLake. Our innovation rests on three novel technical contributions:

Model Attestation Token (MAT): The central artifact of SecureModelGate is a cryptographically signed certificate that binds two distinct identity signals. First, a SHA-256 weight hash confirms file-level integrity. Second, a 128-dimensional behavioral fingerprint is derived by running the model on a fixed canonical test corpus and recording output probability distribution, per-layer activation statistics, and confidence score distribution. This binding is the key novelty: an adversary controlling fine-tuning can embed a backdoor while maintaining a plausible weight hash. By requiring the behavioral signal to match expected patterns, a tampered model cannot produce a valid MAT even when its hash appears legitimate. The MAT is a JWT signed with an HPE-internal RSA-4096 key.

Model Bill of Materials (MBOM): Inspired by SBOM mandates under US Executive Order 14028 [9], we introduce the MBOM - a machine-readable, signed record of base model lineage, fine-tuning provenance, and dependency hashes - directly supporting EU AI Act Article 13 transparency obligations, making SecureModelGate a compliance-enabling tool.

Kubernetes Admission Enforcement: SecureModelGate integrates with HPE Ezmeral via a Kubernetes validating admission webhook. Any pod spec containing a model-serving container must carry a valid MAT annotation. The webhook verifies the token signature and expiry before admitting the pod. Workloads without a valid token are blocked at the scheduler - not merely flagged - ensuring enforcement is mandatory.

The pipeline architecture: every model transits an Intake Gateway orchestrating three parallel modules. The static analysis module applies Kolmogorov-Smirnov tests on weight tensor distributions, inspects the computation graph for anomalous branches, and checks hashes against a known-bad registry. The behavioral fingerprinting module uses architecture-appropriate canonical corpora (image classifiers: 1,000-image ImageNet holdout; text models: 200-prompt fixed set). Only models passing all three checks receive a signed MAT; the K8s webhook then enforces the token before any deployment on HPE Ezmeral proceeds.




Figure 1: SecureModelGate end-to-end model attestation pipeline. Red zone = untrusted ingress; Green zone = attested deployment on HPE Ezmeral.

Evidence the solution works

Evaluation used the NIST TrojAI benchmark [9]: 120 PyTorch classifiers (60 clean, 60 backdoored with patch, blended, and Instagram-filter triggers) on an HPE ProLiant DL380 Gen10 with NVIDIA T4, running HPE Ezmeral 5.5 on RHEL 8.

Table 1: Detection performance on NIST TrojAI benchmark (n=120 models)

Method	Detection Rate	False Positive Rate	Latency
Neural Cleanse [7]	83.3%	7.2%	312 s
STRIP [8]	78.1%	11.5%	28 s
Static only (ours)	80.0%	5.3%	12 s
Behavioral only (ours)	86.7%	6.2%	31 s
SecureModelGate (full)	91.3%	4.1%	38 s

SecureModelGate achieves 91.3% detection at 4.1% FPR, outperforming Neural Cleanse [7] (83.3%/7.2%) and STRIP [8] (78.1%/11.5%). Signal complementarity is key: 12 models evaded static analysis but were caught by behavioral fingerprinting. All 60 backdoored deployment attempts were blocked (100% enforcement); webhook latency 43 ms; MAT issuance 38 s per model (one-time cost).




Figure 2: Detection performance comparison on NIST TrojAI benchmark (n=120 models). SecureModelGate achieves best detection rate (91.3%) and lowest false positive rate (4.1%).

Competitive approaches

Table 2: Competitive landscape - SecureModelGate vs. existing tools

Tool	Behav. Attest.	MBOM	K8s Enforce	Signed Token
Sigstore / cosign [6]	No	No	Yes	Yes
HuggingFace scanning	No	No	No	No
MLflow registry [4]	No	No	No	No
Neural Cleanse [7]	Yes	No	No	No
SecureModelGate (ours)	Yes	Yes	Yes	Yes

Sigstore/cosign [6] signs on file hashes only - no behavioral signal or MBOM. HuggingFace scanning is basic malware detection with no enforcement. MLflow [4] tracks lineage without integrity checks. Neural Cleanse [7] and STRIP [8] are research tools with no enterprise integration. SecureModelGate is the first to combine behavioral attestation, MBOM, Kubernetes enforcement, and a signed token on HPE infrastructure.

Current status

The PoC is complete: static analysis and behavioral fingerprinting (Python/PyTorch), MAT issuance (JWT/RSA-4096), Kubernetes admission webhook on HPE Ezmeral 5.5, and full evaluation on 120 NIST TrojAI models validated end-to-end on HPE ProLiant hardware.

Next steps

SecureModelGate can be productized as a Trusted AI Pipeline module for HPE GreenLake and Ezmeral with no customer workload changes. Target verticals are BFSI, healthcare, and defence - all facing AI governance mandates. The EU AI Act (August 2026) Article 13 MBOM obligations and US NIST AI RMF supply chain requirements are directly addressed. No hyperscaler currently offers behavioral model attestation as a managed service, creating a clear first-mover opportunity for HPE across Security, Ezmeral, and GreenLake teams.

Planned extensions:

- Extend MBOM to capture fine-tuning data lineage automatically via PyTorch training hooks (Q1 2026)
- Integrate MAT status in OpsRamp for fleet-wide attestation observability across GreenLake (Q2 2026)
- Extend to LLM and GGUF-format quantized models used in GreenLake AI inference (Q2 2026)
- Continuous re-attestation: periodic re-fingerprinting of deployed models to detect tampering (Q3 2026)

Acknowledgements

The authors thank the HPE Security and Ezmeral Platform teams for infrastructure support, and acknowledge the NIST TrojAI program for the benchmark dataset.

References

- [1] HuggingFace, "Hub Statistics," 2024. huggingface.co
- [2] X. Chen et al., "Targeted backdoor attacks via data poisoning," arXiv:1712.05526, 2017.
- [3] T. Gu et al., "BadNets: ML model supply chain vulnerabilities," arXiv:1708.06733, 2019.
- [4] M. Zaharia et al., "MLflow," IEEE Data Eng. Bull., vol. 41, no. 4, 2018.
- [5] MITRE, "MITRE ATLAS," 2023. atlas.mitre.org
- [6] L. Dex et al., "Sigstore," Proc. ACM CCS, 2022.
- [7] B. Wang et al., "Neural Cleanse," Proc. IEEE S&P, 2019.
- [8] Y. Gao et al., "STRIP," Proc. ACSAC, 2019.
- [9] The White House, "EO 14028 - Improving Nation's Cybersecurity," 2021.
- [10] European Parliament, "EU AI Act (Reg. 2024/1689)," 2024.
- [11] IARPA, "NIST TrojAI Benchmark," 2021. pages.nist.gov/trojai